

# VeriTracer: Context-enriched tracer for floating-point arithmetic analysis

ARITH 25, Amherst MA USA, June 2018

---

Yohan Chatelain<sup>1,5</sup>

Pablo de Oliveira Castro<sup>1,5</sup>

Eric Petit<sup>2,5</sup>

David Defour<sup>3</sup>

Jordan Bieder<sup>4,5</sup>

Marc Torrent<sup>4,5</sup>



<sup>1</sup>Université de Versailles Saint-Quentin-en-Yvelines (UVSQ)

<sup>2</sup>Intel

<sup>3</sup>Université de Perpignan Via Domitia (UPVD)

<sup>4</sup>CEA, DAM, DIF

<sup>5</sup>Exascale Computing Research (ECR)

# Table of contents

1. Introduction

2. Veritracer

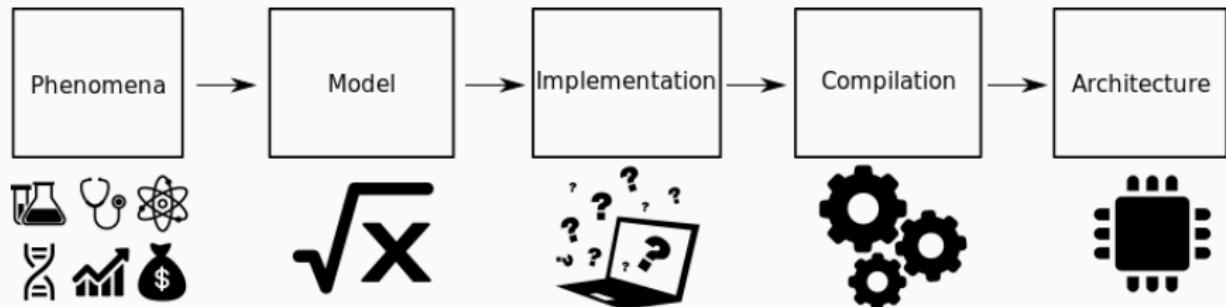
3. Experiments

4. Conclusion

# Introduction

---

Building fast and robust applications for HPC is a complex task !



- At each step, numerical bugs can be introduced

**Objective:** Tools to track and analyze numerical bugs

# Existing tools

	Method	Name	Implementation
Debugging	Stochastic Arithmetic	CADNA [9] VERROU [5] Verifierlo [4]	CESTAC/DSA (library) CESTAC (Valgrind) MCA (LLVM)
	Extended Precision (EP)	HPC Craft [10] FpDebug [3] Herbgrind [15]	Exponent comparison (DynInst) MPFR (Valgrind) MPFR (Valgrind)

	Name	Method	Mixed-Precision	Any-Precision
Optimization	Verifierlo [4]	MCA (User specific)	✓	✓
	HPC Craft [10]	Bitmask (ref value)	✓	✓
	Promise [8]	CESTAC/DSA ( $\Delta - \text{debug}$ )	✓	
	Precimonious [14]	EP ( $\Delta - \text{debug}$ )	✓	
	Herbie [12]	EP (Rewriting)		

# Muller's sequence[2] overview

exact solution

$$u_n = \left\{ \begin{array}{l} u_0 = 2 \\ u_1 = -4 \\ u_{n+1} = 111 - \frac{1130}{u_n} + \frac{3000}{u_n u_{n-1}} \end{array} \right\} \quad \lim_{n \rightarrow \infty} u_n =$$

6

finite precision

100

CADNA

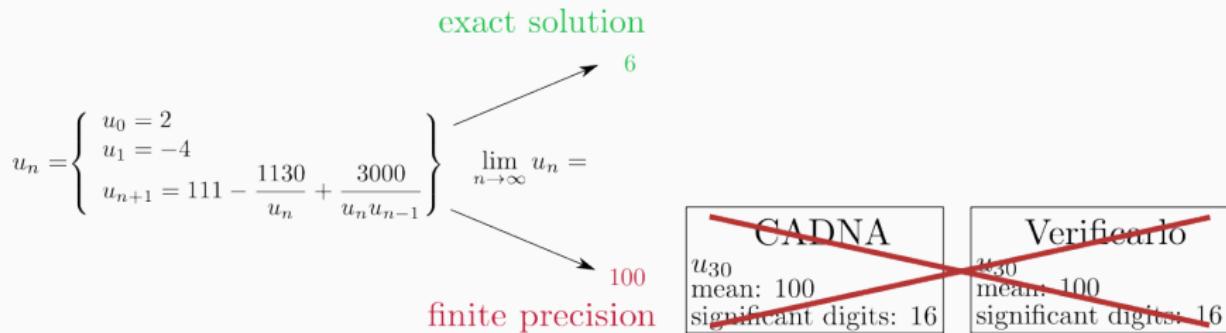
$u_{30}$   
mean: 100  
significant digits: 16

Verificarlo

$u_{30}$   
mean: 100  
significant digits: 16

Stochastic methods execute several samples by introducing random errors

# Muller's sequence[2] overview



## Without an exact solution

- Only checking the final result may lead to wrong conclusions !

## Current challenges

---

Existing tools explore the **spatial** dimension of numerical computations:

- which variable or operation is imprecise
- which function can be switched into lower precision
- but programs have different numerical requirements over time

⇒ Need to explore the **temporal** dimension

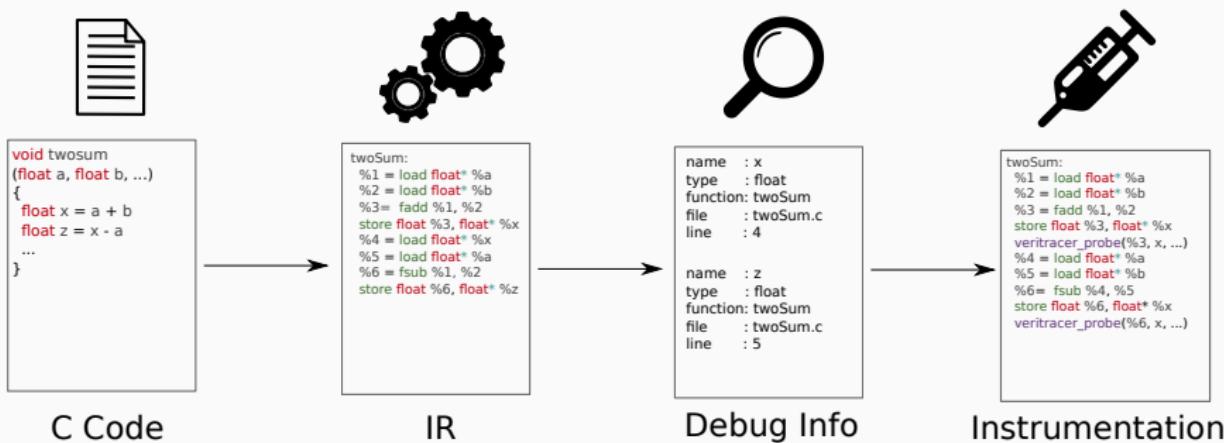
# Veritacer

---

# Instrumentation

## Veritracer LLVM pass

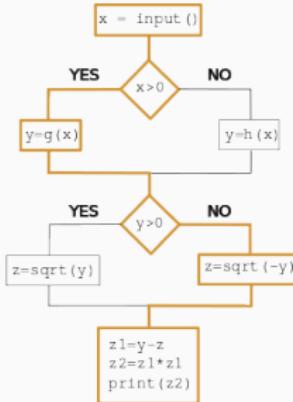
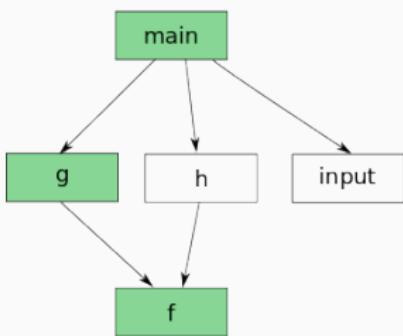
- Accepts any source code accepted by the LLVM frontend
- Transforms into an LLVM Intermediate Representation
- Searches debug information through LLVM IR
- Instruments the IR code with veritracer probes



# Information Flow Analysis

## Ansychronous analysis

- Veritracer works in an asynchronous way
- Information flow is necessary for gathering the traces
- Three levels of granularities are analyzed:
  - Context Analysis: considers the context of a function call
  - Control-Flow Analysis: considers the flow path taken
  - Data-Flow Analysis: considers the dependencies between variables



a	=	b	+	c
d	=	e	-	f
g	=	a	*	f
h	=	a	+	d
i	=	g	/	h

# Pinpointing relevant function

## Narrowing the search space

To narrow the search space, we:

**Scan** functions with FP values

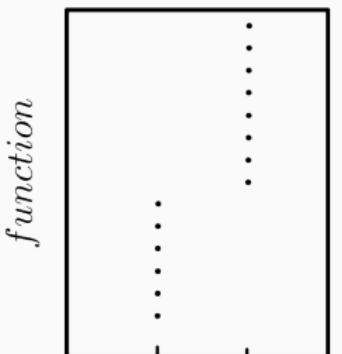
**Narrow** the set by removing non-contributing functions

$$f_{\delta e_{\max}}(x) \leq \epsilon,$$

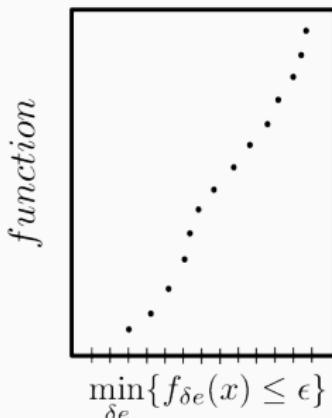
**Order** functions by *numerical stress resistance*

$$\min_{\delta e} \{ f_{\delta e}(x) \leq \epsilon \}$$

*Narrowing*



*Ordering*



$$\min_{\delta e} \{ f_{\delta e}(x) \leq \epsilon \}$$

# Implementation

## Extends verificarlo to

- Trace the precision of FP variables over time
- Provide contextual information on traced variables
  - Current version implements Context Analysis
  - Inserts call to `backtrace()` (GNU C library)
  - Analyzes backtraces to detect flow divergences
- All in an automatic way

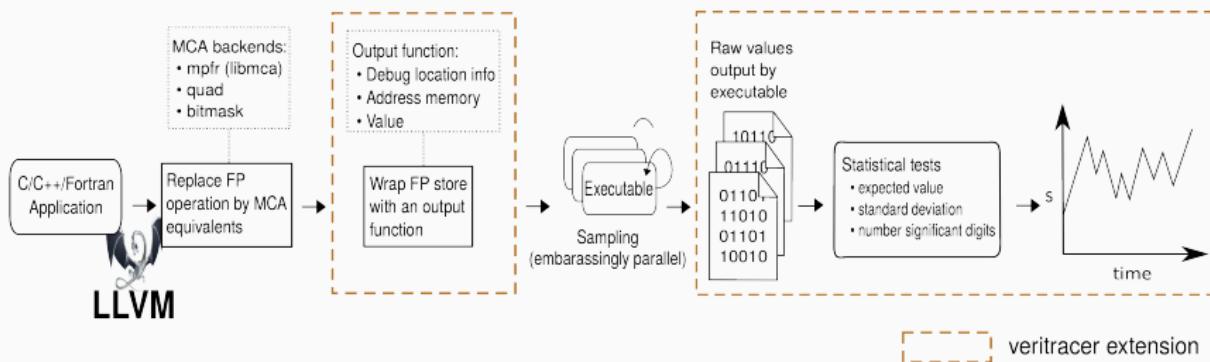


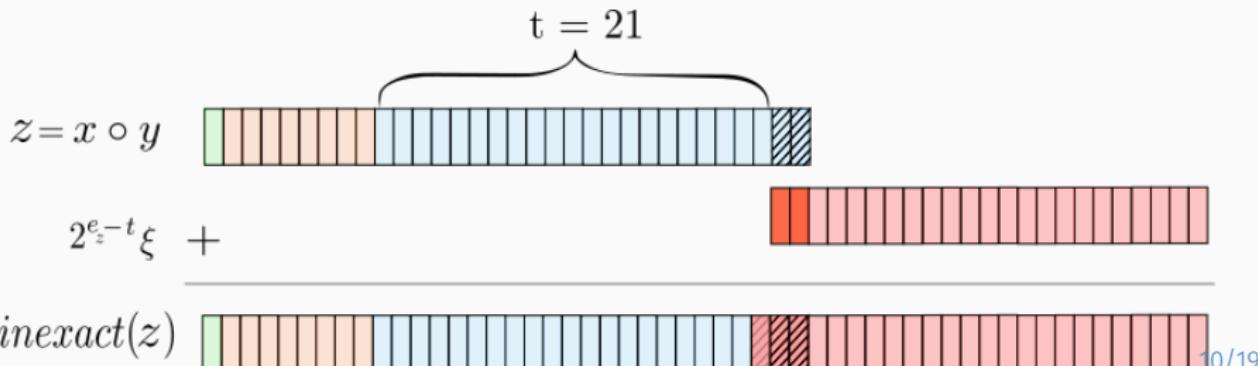
Figure 1: Veritracer's workflow

$$\text{inexact}(x) = x + \beta^{e_x-t} \xi$$

- $e_x$  is the magnitude of  $x$
- $t$  the virtual precision
- $\xi$  a random variable between  $[-\frac{1}{2}, \frac{1}{2}]$

## Random Rounding mode

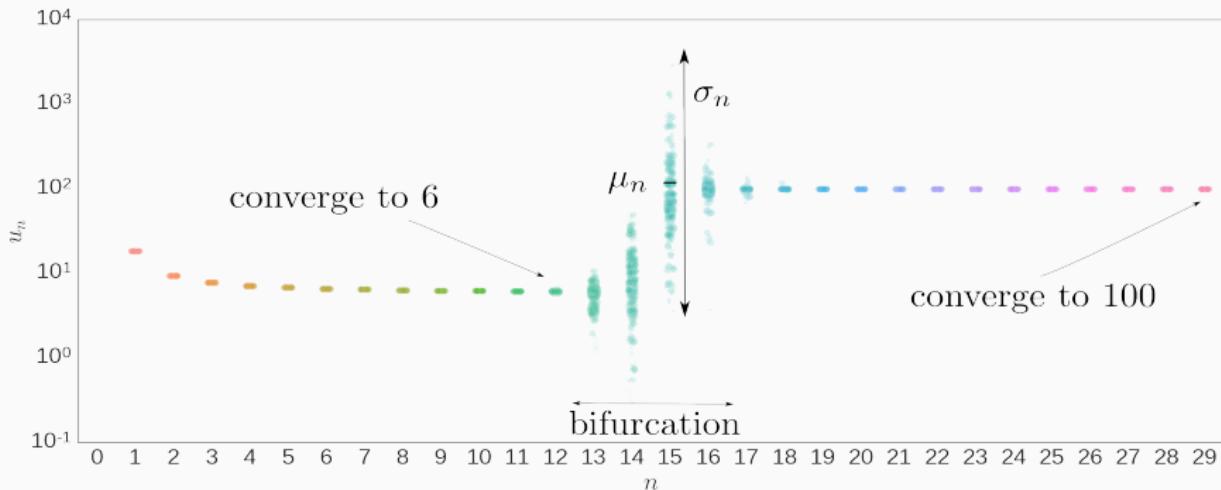
- $RR(x \circ y) = \text{inexact}(x \circ y), \circ \in \{+, -, \times, \setminus\}$



## Experiments

---

# Muller's sequence with veritracer

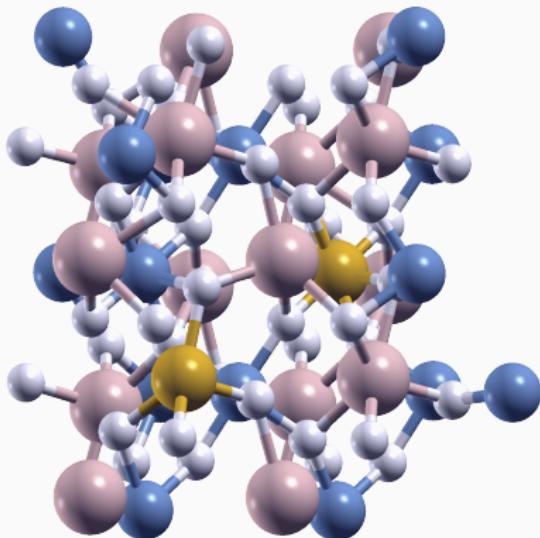


**Figure 2:** Muller's sequence evaluation over 500 samples with veritracer. The sequence converges to 6 before bifurcating and then converges to 100.

- Estimate the number of significant digits by using N samples
$$\tilde{s} = -\log_{10} \left( \frac{\tilde{\sigma}}{\tilde{\mu}} \right)$$
- $\tilde{\mu}$ : empirical mean and  $\tilde{\sigma}$ : empirical standard deviation

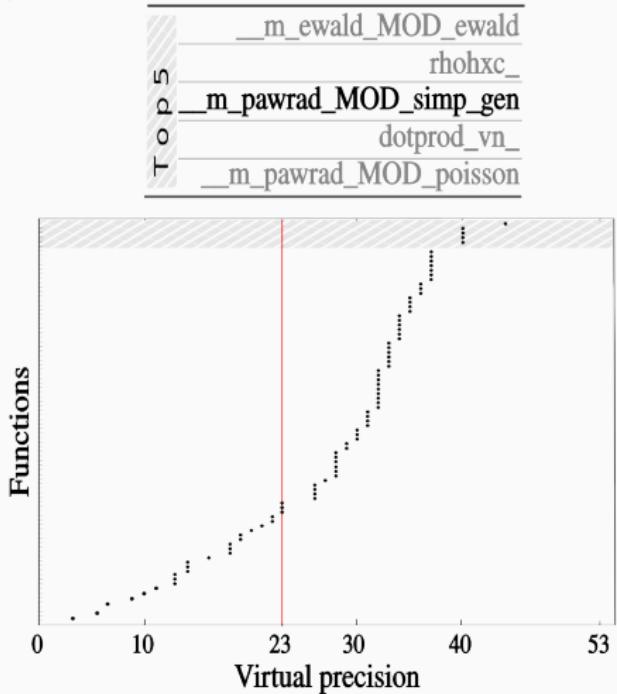
## ABINIT [7]

- Calculates observable properties of materials (optical, mechanical, vibrational)
- Works on any chemical composition (molecules, nanostructures, solids)



**Figure 3:** Sound velocity calculation in an earth mantle component ( $MgSiO_3$  perovskite with Al impurities) [1]

# Ordering the search space



## Functions below 23bits

- are more resistant
- can potentially be changed to single precision

## Functions above 23bits

- are less resistant
- require further analyses

## Interesting function among the top 5

`__m_pawrad_MOD_simp_gen`

# Simp\_gen: Evaluation with Veritracer

Computes an integral by Simpsons' rule over a generalized 1D-grid

- Can be seen as a dot product

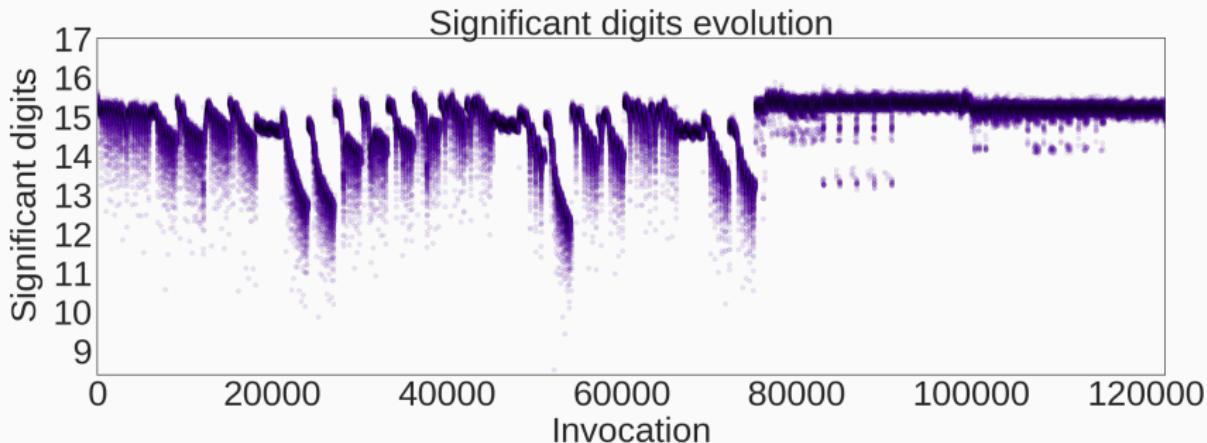
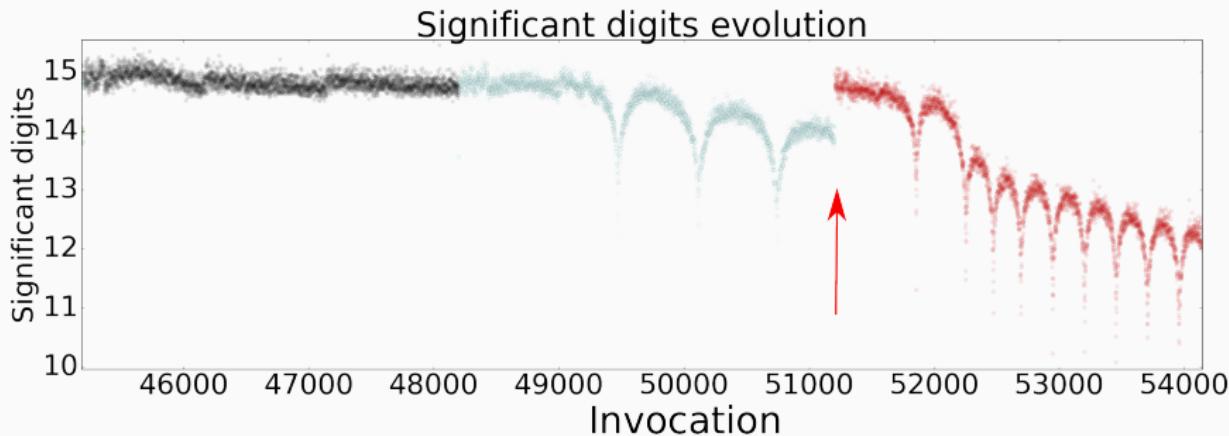


Figure 4: Evolution of the value returned by simp\_gen with  $t = 53$  in Random Rounding mode with 24 samples in parallel on the CINES Occigen cluster (2106 nodes x2 Intel 12-Cores (E5-2690V3@2.6 GHz))

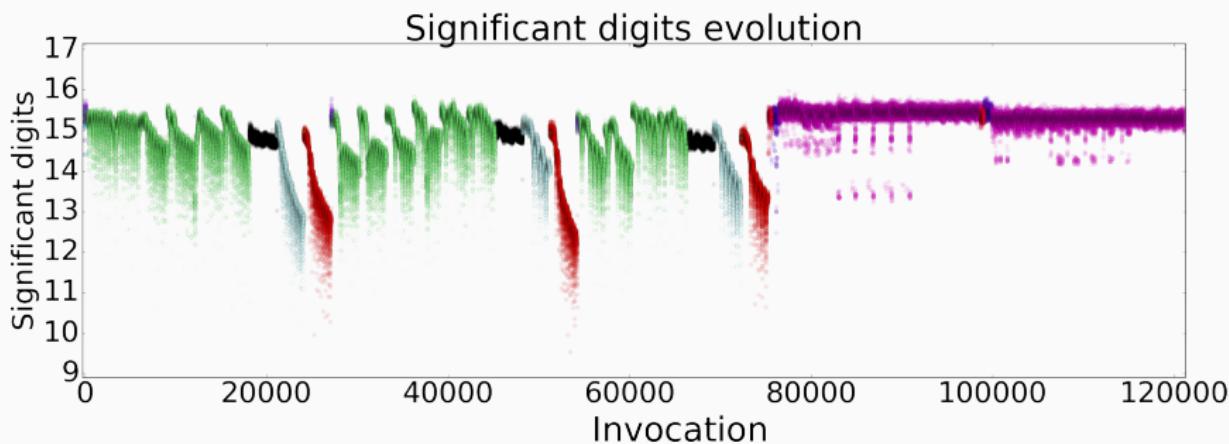
## Simp\_gen: the importance of the context analysis

- Colors represent the different CSPs
- Red arrow seems to point an enhancement of the quality ...
- ... but red points come from another context

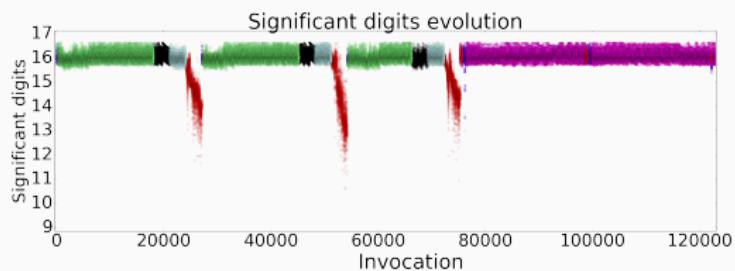
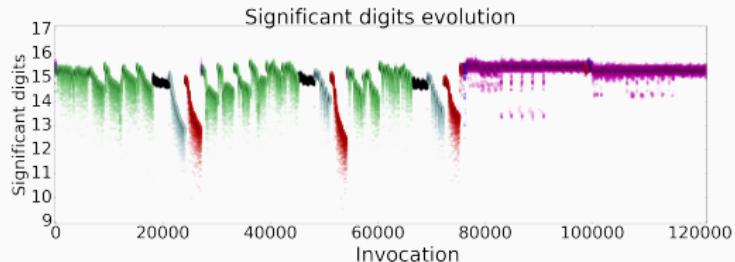


## Simp\_gen: the importance of the context analysis

- Colors represent the different CSPs
- Red arrow seems to point an enhancement of the quality ...
- ... but red points come from another context
- Hence the importance of separating contexts



# Simp\_gen: A compensated version



## Simp\_gen: Altered version

- Can be seen a dot product
- Replaces by a compensated version Dot2 [11]
- Implemented with libeft [6]
- The precision of 30/31 CSPs is improved
- 1 CSP has a low precision due to reentrance of the error

## Conclusion

---

# Conclusion

Veritracer is a numerical debugger/optimizer which

- Automatically instruments codes through the verificarlo framework
- Automatically plots the temporal numerical quality of variables
- Provides contextual information

Veritracer can be used on real world HPC applications to

- Detect the numerical sensitivity of functions
- Classify functions according to their numerical sensitivity
- Observe the numerical behavior of functions over time

Thank you for your attention



Available on github:

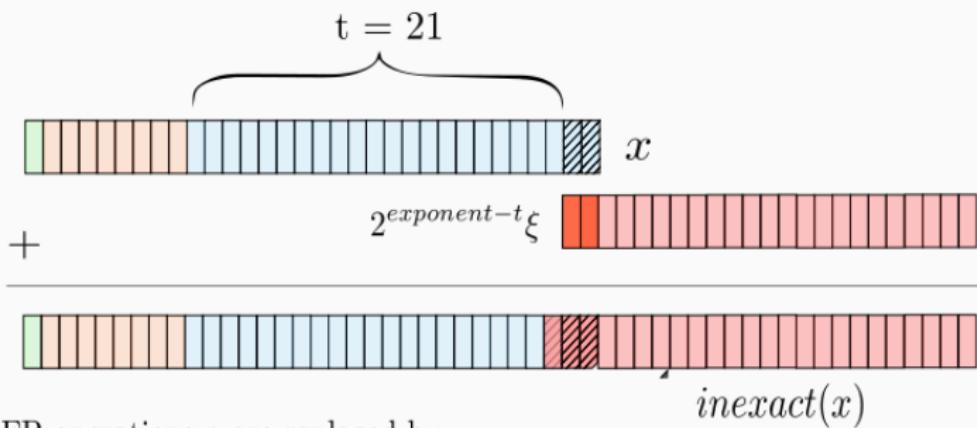
[github.com/verificarlo/verificarlo](https://github.com/verificarlo/verificarlo)

[github.com/verificarlo/verificarlo/tree/veritracer](https://github.com/verificarlo/verificarlo/tree/veritracer)

## Random Rounding mode (MCA)

### Monte Carlo Arithmetic [parker1997monte]

$$\text{inexact}(x) = x + 2^{\text{exponent}-t} \xi, \quad \xi \in [-\frac{1}{2}, \frac{1}{2}], \quad t \text{ virtual precision}$$



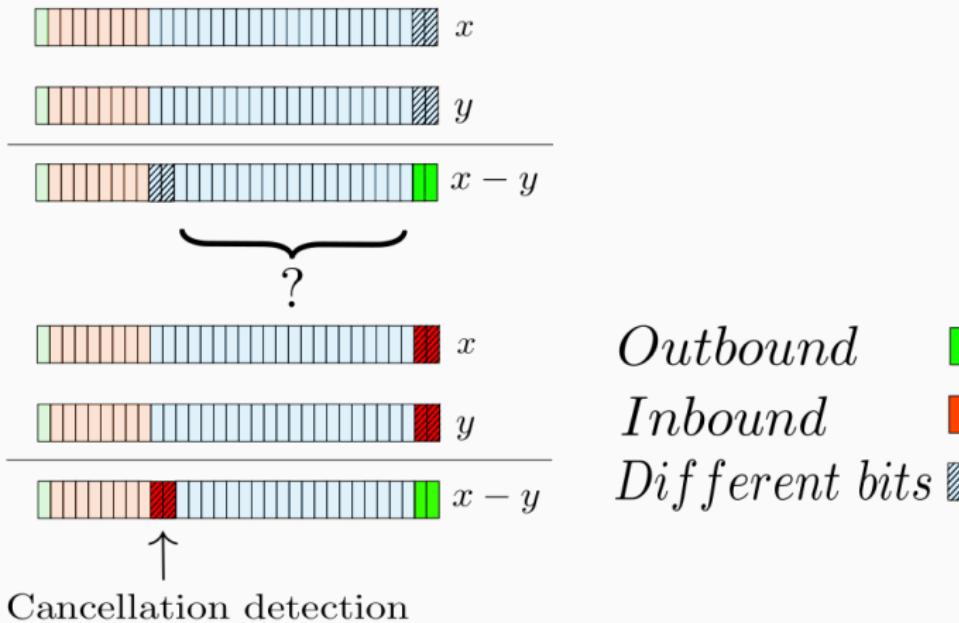
FP operations  $\circ$  are replaced by:

$$RR(x \circ y) = \text{round}(\text{inexact}(x \circ y))$$

# Full MCA mode

$$x \sim y$$

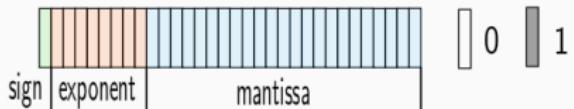
$$mca(x) = \text{round}(\text{inexact}(\text{inexact}(x) - \text{inexact}(y)))$$



# Bitmask backend

VERIFICARLO\_PRECISION=11

IEEE754 Single precision 32-bit



Mode ZERO



AND



3.1415927

Mask

3.140625

Mode INV



Mode RAND



OR



Mask

XOR



Mask

3.1407475

# Errors propagation

```
subroutine pawpssp_calc(...)  
...  
if (testval) then  
...  
    call simp_gen(qq, nhat, vale_mesh)  
    qq=zion/four_pi-qq  
end if  
  
call atompaw_shpfun(..., intg, nhat)  
  
nhat(1:msz)=qq*nhat(1:msz)  
tnvale(1:msz)=tnvale(1:msz)+nhat(1:msz)  
...  
call pawpssp_cg(..., tnvale, ...)
```

♠

```
subroutine pawpssp_cg(..., nr, ...)  
...  
do ir=1,mesh_size  
    rnr(ir)=radmesh%rad(ir)*nr(ir)  
end do  
  
do ir=1,mesh_size  
    if (abs(rnr(ir))>1.d-20)  
        ff(ir)=sin(arg*radmesh%rad(ir))*rnr(ir)  
    end do  
...  
call simp_gen(r1torm, ff, radmesh)
```

```
subroutine atompaw_shpfun(..., norm, shapefunc)  
...  
simp_gen(norm,r2k,mesh)  
norm=one/norm  
shapefunc(1:ishp)=shapefunc(1:ishp)*norm
```

inout

out

out

```
subroutine simp_gen(intg, func, radmesh)  
...  
nn=radmesh%int_meshsz  
simp=zero  
do i=1,nn  
    simp=simp+func(i)*radmesh%simfact(i)  
end do  
...  
intg=simp+resid  
end subroutine
```

# Veritracer: the LLVM IR

```
void twoSum
(double a, double b, double *x_ptr, double *err_ptr) {
    double x = a + b;
    double z = x - a;
    double e = (a - (x - z)) + (b - z);
    *x_ptr = x;
    *err_ptr = e;
}
```

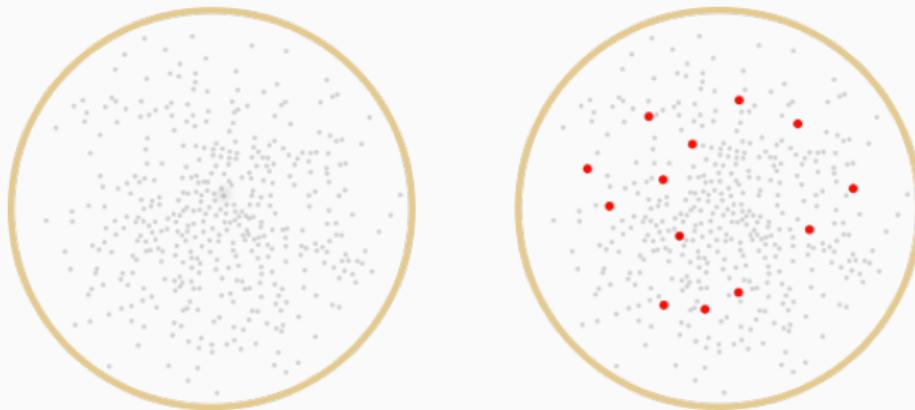
```
define void @twoSum
(float %a, float %b, float* %x_ptr, float* %e_ptr)
{
%5 = load float* %a
%6 = load float* %b
%7 = fadd float %5, %6 ; a+b
store float %7, float* %x ; x = a+b
%8 = load float* %x
%9 = load float* %a
%10 = fsub float %8, %9 ; x-a
store float %10, float* %z ; z = x-a
%11 = load float* %a
%12 = load float* %x
%13 = load float* %z
%14 = fsub float %12, %13 ; x-z
%15 = fsub float %11, %14 ; a-(x-z)
%16 = load float* %b
%17 = load float* %z
%18 = fsub float %16, %17 ; b-z
%19 = fadd float %15, %18 ; a-(x-z)+(b-z)
store float %19, float* %e ; e = a-(x-z)+(b-z)
%20 = load float* %x
%21 = load float** %3
store float %20, float* %21 ; *x_ptr = x
%22 = load float* %e
%23 = load float** %4,
store float %22, float* %23 ; *e_ptr = e
ret void
}
```

```
define void @twoSum
(float %a, float %b, float* %x_ptr, float* %err_ptr)
{
%1 = fadd float %a, %b ; a+b      (x)
%2 = fsub float %1, %a ; (a+b)-a (z)
%3 = fsub float %1, %2 ; x-z
%4 = fsub float %a, %3 ; a-(x-z)
%5 = fsub float %b, %2 ; b-z
%6 = fadd float %5, %4 ; (b-z)+(a-(x-z))
store float %1, float* %x_ptr ; *x_ptr = x
store float %6, float* %err_ptr ; *err_ptr = e
ret void
}
```

# Reducing the search space

## Instrumenting one function of the program at a time

- Executes each version with the maximal MCA error
- Compares the MCA computed value with the reference:
  - Different : the function is tagged as critical
  - Identical : the function is tagged not critical



2952 function  $\Rightarrow$  88 functions  
Reduces the search space by  $\times 30$

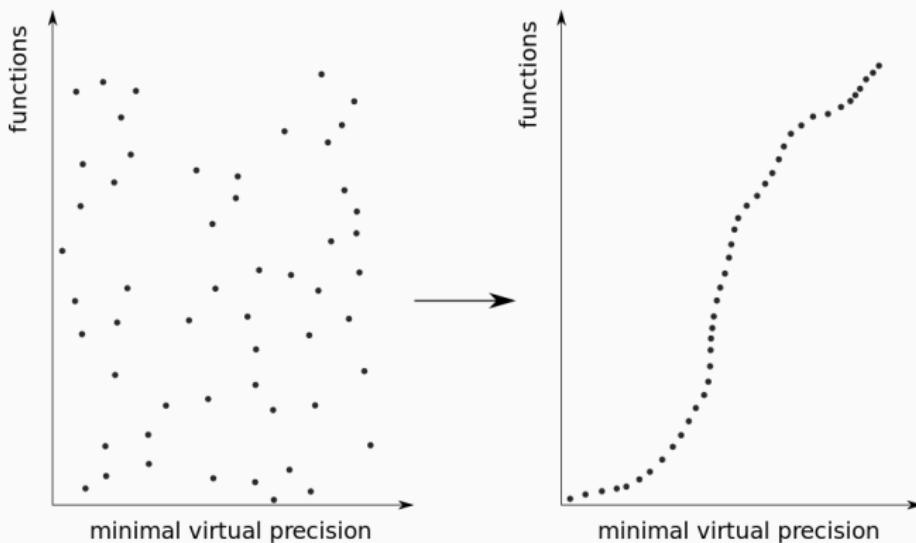
# Ordering the search space

## Problem

Functions do not have the same numerical behavior: some functions are less sensitive to numerical errors than others

## Idea

Classifying functions according to their numerical stress resistance



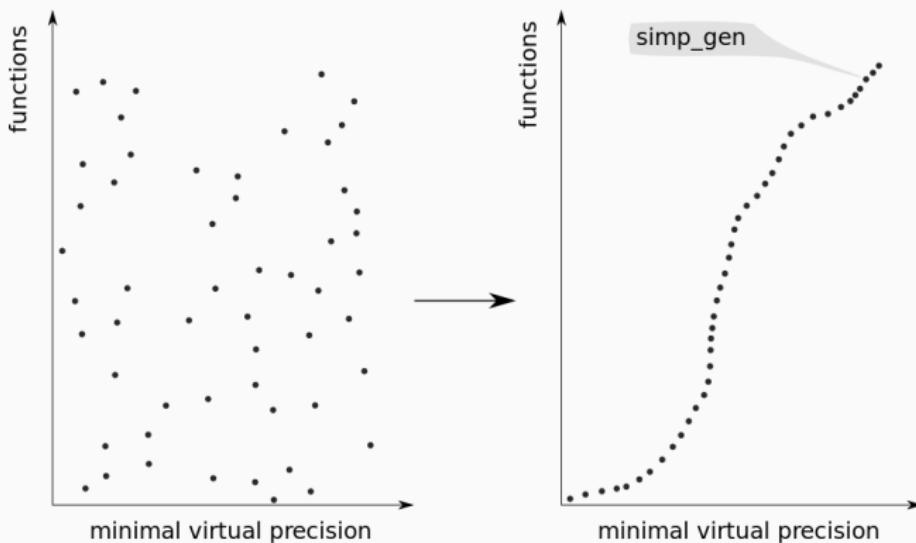
# Ordering the search space

## Problem

Functions do not have the same numerical behavior: some functions are less sensitive to numerical errors than others

## Idea

Classifying functions according to their numerical stress resistance



## References i

- [1] ABINIT.  
**Silicon carbide, 2016.**  
[Online; accessed May 15, 2018].
- [2] J.-C. Bajard, D. Michelucci, J.-M. Moreau, and J.-M. Muller.  
**Introduction to the Special Issue "Real Numbers and Computers".**  
In *The Journal of Universal Computer Science*, pages 436–438.  
Springer, 1996.
- [3] F. Benz, A. Hildebrandt, and S. Hack.  
**A dynamic program analysis to find floating-point accuracy problems.**  
*ACM SIGPLAN*, 47(6):453–462, 2012.

## References ii

- [4] C. Denis, P. de Oliveira Castro, and E. Petit.  
**Verificarlo: checking floating point accuracy through monte carlo arithmetic.**  
In *Computer Arithmetic (ARITH), 23nd Symposium on*, pages 55–62.  
IEEE, 2016.
- [5] F. Févotte and B. Lathuiliere.  
**VERROU: a CESTAC evaluation without recompilation.**  
*SCAN 2016*, page 47, 2016.
- [6] F. Févotte and B. Lathuilière.  
**LibEFT: a library implementing Error-Free transformations**, 2017.
- [7] X. Gonze, F. Jollet, et al.  
**Recent developments in the ABINIT software package.**  
*Computer Physics Communications*, 205:106–131, 2016.

## References iii

- [8] S. Graillat, F. Jézéquel, R. Picot, F. Févotte, and B. Lathuilière.  
**Auto-tuning for floating-point precision with Discrete Stochastic Arithmetic.**  
working paper or preprint, June 2016.
- [9] F. Jézéquel and J.-M. Chesneaux.  
**CADNA: a library for estimating round-off error propagation.**  
*Computer Physics Communications*, 178(12), 2008.
- [10] M. O. Lam, J. K. Hollingsworth, and G. Stewart.  
**Dynamic floating-point cancellation detection.**  
*Parallel Computing*, 39(3):146–155, 2013.
- [11] T. Ogita, S. M. Rump, and S. Oishi.  
**Accurate sum and dot product.**  
*SIAM Journal on Scientific Computing*, 26(6):1955–1988, 2005.

## References iv

- [12] P. Panchekha, A. Sanchez-Stern, J. R. Wilcox, and Z. Tatlock.  
**Automatically improving accuracy for floating point expressions.**  
In *ACM SIGPLAN Notices*, volume 50, pages 1–11. ACM, 2015.
- [13] D. S. Parker.  
**Monte Carlo Arithmetic: exploiting randomness in floating-point arithmetic.**  
University of California. Computer Science Department, 1997.
- [14] C. Rubio-González, C. Nguyen, H. D. Nguyen, J. Demmel, W. Kahan,  
K. Sen, D. H. Bailey, C. Iancu, and D. Hough.  
**Precimonious: Tuning assistant for floating-point precision.**  
In *Proceedings of the International Conference on High Performance Computing, Networking, Storage and Analysis*, page 27. ACM, 2013.
- [15] A. Sanchez-Stern, P. Panchekha, S. Lerner, and Z. Tatlock.  
**Finding root causes of floating point error with herbgrind.**  
*arXiv preprint arXiv:1705.10416*, 2017.